

Intro to Least Squares

28 May 2026 08:01



06-intro-to
-least-squ...

Introduction to Least Squares

Purpose

of what, to what?

Understand least squares as finding the best approximate fit when an exact solution is not possible.

This notebook uses a small line-fitting example to connect three ideas:

- residuals as prediction errors
- least squares as minimising the sum of squared residuals
- the matrix form used by `np.linalg.lstsq`

We will keep the example small and visual.

What will I know at end?

Imports and helpers

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

def setup_plot(title, xlim=(0.5, 4.5), ylim=(1.5, 4.25), figsize=(5.5, 4)):
    fig, ax = plt.subplots(figsize=figsize)
    ax.set_xlim(*xlim)
    ax.set_ylim(*ylim)
    ax.grid(True, alpha=0.3)
    ax.set_xlabel("x")
    ax.set_ylabel("y")
    ax.set_title(title)
    return fig, ax

def predict_line(x_values, slope, intercept):
    return slope * x_values + intercept

def sum_squared_errors(actual, predicted):
    residuals = actual - predicted
    return np.sum(residuals ** 2)
```

```
def sum_squared_errors(actual, predicted):  
    residuals = actual - predicted  
    return np.sum(residuals ** 2)
```

A simple fitting problem

Suppose we have some points:


$(1, 2), (2, 2.5), (3, 3.5), (4, 3.8)$

They do not lie perfectly on a straight line.

But we can still ask:

what line fits these points reasonably well?

```
In [2]: x = np.array([1, 2, 3, 4])  
        y = np.array([2.0, 2.5, 3.5, 3.8])  
  
        fig, ax = setup_plot("Data points")  
        ax.scatter(x, y, label="actual points")  
        ax.legend()  
        plt.show()
```



A line model

A straight line has the form:

A line model

A straight line has the form:

$$y = mx + c$$

where m is the slope and c is the intercept.

For a given m and c , we can predict a y -value for every x -value.

```
In [3]: first_m = 0.7
first_c = 1.2


first_predicted_y = predict_line(x, first_m, first_c)

print("actual y: ", y)
print("predicted y:", first_predicted_y)
```

```
actual y: [2. 2.5 3.5 3.8]
predicted y: [1.9 2.6 3.3 4. ]
```

```
In [4]: line_x = np.linspace(0.5, 4.5, 100)
line_y = predict_line(line_x, first_m, first_c)
```

```
plt.show()
```



Residuals

A residual is the difference between an actual value and a predicted value.

$$\text{residual} = y - \hat{y}$$

Here, $\hat{y}$ is the predicted value.

→ Because we are projecting onto a 1-D column space.

candidate
line
of
first

Bad fit

residuals

1.7
2

$$\text{residual} = y - \hat{y}$$

Here, $\hat{y}$ is the predicted value.

These are vertical residuals because this model predicts a y-value from each x-value.

```
In [5]: first_residuals = y - first_predicted_y
print("residuals:", first_residuals)
```


```
residuals: [ 0.1 -0.1  0.2 -0.2]
```

```
In [6]: fig, ax = setup_plot("Residuals as vertical errors")
ax.scatter(x, y, label="actual points")
ax.plot(line_x, line_y, label="candidate line")


for xi, actual, predicted in zip(x, y, first_predicted_y):
    ax.plot([xi, xi], [predicted, actual], linestyle="--", linewidth=1.5)

ax.legend()
plt.show()
```

there other approaches
th different definitions
of 'best'?



Are the with
 Do we need a
 'obviously we cant',
 kelly millions' ?
 Is that the
 only reason?
 Any downsides?



The dashed lines show the vertical errors between the candidate line and the data points.

Least squares gives a precise meaning to the word **best**:

the best line is the line with the smallest sum of squared residuals.

Why square the residuals?

Residuals can be positive or negative.

If we simply add them, they can cancel out.

So least squares uses the sum of squared residuals:

$$\sum_i (y_i - \hat{y}_i)^2$$

Squaring does two useful things:

- negative and positive errors both become positive
- larger errors are penalised more heavily

```
In [7]: first_ssr = sum_squared_errors(y, first_predicted_y)
print("candidate line sum of squared residuals:", first_ssr)
candidate line sum of squared residuals: 0.10000000000000009
```

Let NumPy find the least-squares line

We will use `np.polyfit` first as a convenient way to find the best line.

Then we will look at the matrix version of the same idea.

```
In [8]: best_m, best_c = np.polyfit(x, y, deg=1)
print("best slope: ", best_m)
print("best intercept:", best_c)
best slope: 0.6399999999999998
best intercept: 1.3500000000000008

In [9]: best_predicted_y = predict_line(x, best_m, best_c)
best_residuals = y - best_predicted_y
best_ssr = sum_squared_errors(y, best_predicted_y)
```

```

best_residuals = y - best_predicted_y
best_ssr = sum_squared_errors(y, best_predicted_y)

print("best predicted y:", best_predicted_y)
print("best residuals: ", best_residuals)
print("best sum of squared residuals:", best_ssr)
print("candidate line sum of squared residuals:", first_ssr)

```

```

best predicted y: [1.99 2.63 3.27 3.91]
best residuals:  [ 0.01 -0.13  0.23 -0.11]
best sum of squared residuals: 0.08199999999999993
candidate line sum of squared residuals: 0.10000000000000009

```

```


In [10]: best_line_y = predict_line(line_x, best_m, best_c)

fig, ax = setup_plot("Least-squares fit")
ax.scatter(x, y, label="actual points")
ax.plot(line_x, best_line_y, label="least-squares line")

for xi, actual, predicted in zip(x, y, best_predicted_y):
    ax.plot([xi, xi], [predicted, actual], linestyle="--", linewidth=1.5)

ax.legend()
plt.show()

```



The least-squares line has a smaller sum of squared residuals than the earlier candidate

The least-squares line has a smaller sum of squared residuals than the earlier candidate line.

That is what **best fit** means in this notebook.

Best fit has the best of Any line

Matrix view of the same problem

The line equation is:

$$\hat{y} = mx + c$$

For one data point x_i , the prediction is:

$$\hat{y}_i = mx_i + c$$

We can write the same calculation as a dot product:

$$\hat{y}_i = [x_i \ 1] \begin{bmatrix} m \\ c \end{bmatrix}$$

because:

$$[x_i \ 1] \begin{bmatrix} m \\ c \end{bmatrix} = x_i m + 1c = mx_i + c$$



The 1 is there because the intercept c is added once for every prediction.

So each row of the matrix is one small prediction rule:

$$[x_i \ 1]$$

For the x-values 1, 2, 3, 4, the four prediction equations are:

$$\hat{y}_1 = 1m + 1c$$

$$\hat{y}_2 = 2m + 1c$$

$$\hat{y}_3 = 3m + 1c$$

$$\hat{y}_4 = 4m + 1c$$

1, 2, 3, 4
because these are the x values in the given points.

Stacking those equations gives the matrix form:

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \hat{y}_3 \\ \hat{y}_4 \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 2 & 1 \\ 3 & 1 \\ 4 & 1 \end{bmatrix} \begin{bmatrix} m \\ c \end{bmatrix}$$

More compactly:

$$\hat{\mathbf{y}} = X\boldsymbol{\beta}$$

where:

Maybe we need non integer x example?

$$X = \begin{bmatrix} 1 & 1 \\ 2 & 1 \\ 3 & 1 \\ 4 & 1 \end{bmatrix} \quad \beta = \begin{bmatrix} m \\ c \end{bmatrix}$$

The matrix X is called the design matrix.

important

In this simple line-fitting problem:

Always →

- the first column contains the x-values
- the second column contains ones so that the intercept c is included in every prediction
- the parameter vector contains the unknowns we are trying to learn: m and c

The key idea is that $X\beta$ gives all the predicted y-values at once.

```
In [11]: # We have four observed x-values.
x = np.array([1, 2, 3, 4])

# For a straight line model:
#
# predicted_y = m*x + c
#
# each row of the design matrix needs two values:
#
# [x_value, 1]
#
# The x_value is multiplied by m.
# The 1 is multiplied by c, because the intercept is added once for every prediction.

# np.ones_like(x) creates an array of ones with the same shape as x:
#
# [1, 1, 1, 1]
#
# These ones form the intercept column.
intercept_column = np.ones_like(x)

# np.column_stack(...) puts arrays side by side as columns.
#
# So this combined array:
#
# x          -> [1, 2, 3, 4]
# intercept_column -> [1, 1, 1, 1]
#
# into:
#
# [1, 1],
# [2, 1],
# [3, 1],
# [4, 1]
X = np.column_stack([x, intercept_column])
print(X)
```

```
print(X)
```

```
[[1 1]
 [2 1]
 [3 1]
 [4 1]]
```

Each row of `X` corresponds to one x -value.

For example, the row `[3, 1]` means:

$$\hat{y} = 3m + 1c$$

which is just the original line equation evaluated at $x = 3$.

The unknown parameters are:

$$\beta = \begin{bmatrix} m \\ c \end{bmatrix}$$

Then the predictions are:

$$X\beta$$

```
In [12]: beta = np.array([best_m, best_c])

predictions_from_matrix = X @ beta

print("X @ beta:      ", predictions_from_matrix)
print("best predicted y:", best_predicted_y)
```

```
X @ beta:      [1.99  2.63  3.27  3.91]
best predicted y: [1.99  2.63  3.27  3.91]
```

Why this is not usually an exact solve

There are four data points, so the line would need to satisfy four equations:

$$m \cdot 1 + c \approx 2.0$$

$$m \cdot 2 + c \approx 2.5$$


$$m \cdot 3 + c \approx 3.5$$

$$m \cdot 4 + c \approx 3.8$$

But there are only two unknowns: m and c .

Because the points do not lie perfectly on a straight line, there is usually no exact solution.

Do we need to come here to identify that there are no solutions?



Because the points do not lie perfectly on a straight line, there is usually no exact solution.

So instead of solving $X\beta = y$ exactly, least squares finds the β that gets as close as possible.

Solving with `np.linalg.lstsq`

NumPy can solve the matrix least-squares problem directly:

$$X\beta \approx y$$

This means:

find the best parameter vector β when $X\beta = y$ cannot usually be satisfied exactly.

```
In [13]: beta_lstsq, residual_sums, rank, singular_values = np.linalg.lstsq(X, y, rcond=None)
print("beta from lstsq:", beta_lstsq)
print("slope:           ", beta_lstsq[0])
print("intercept:        ", beta_lstsq[1])
```

```
beta from lstsq: [0.64 1.35]
slope:           0.6400000000000001
intercept:       1.3499999999999994
```

The result matches the line fitted earlier.

~~This~~ is the least-squares solution written in matrix form.

Explain
all?

beta from
lstsq?

Connection back to projection

In the projection notebook, we split a vector into:

parallel component + perpendicular component

Least squares has a similar flavour.

We are not going to fully unpack that geometry here, but the important idea is this:

least squares finds the closest achievable prediction.

The model can only produce predictions of the form $X\beta$.

Least squares chooses the achievable prediction closest to the actual values, leaving a residual error that the model cannot explain.

residual error that the model cannot explain.

Summary

Least squares is useful when an exact fit is not possible.

In this notebook, we fitted a line to a small set of points.

A residual is the prediction error:

$$\text{residual} = y - \hat{y}$$

Least squares chooses the line that minimises:

$$\sum_i (y_i - \hat{y}_i)^2$$

The key takeaways are:

- residuals are the differences between actual and predicted values

- least squares minimises the sum of squared residuals
- a line $y = mx + c$ can be written as a matrix problem
- `np.linalg.lstsq` solves the matrix least-squares problem directly
- geometrically, least squares is connected to finding the closest achievable prediction

The main idea is:

when exact fitting is impossible, least squares gives us a principled best approximate fit.

